

Kaustubh Gupta

kaustubhg73@gmail.com | 8982255359 | [linkedin/kaustubh-gupta-207791223](https://www.linkedin.com/in/kaustubh-gupta-207791223) | scholar | [github/Kaustubh73](https://github.com/Kaustubh73) | [huggingface/kaustubhg73](https://huggingface.co/kaustubhg73)

RESEARCH INTERESTS

Multilingual trustworthy AI; task-conditioned over-refusal in aligned LLMs; mechanistic interpretability; agency-preserving alignment

EDUCATION

BTech in Artificial Intelligence and Engineering Science - 9.34

Hyderabad, IN | 2025

INDIAN INSTITUTE OF TECHNOLOGY HYDERABAD [HIGHEST CGPA IN DEPARTMENT]

Teaching Assistant: Convex Optimization (Jan-May 2024, Jan-May 2025)

Coursework: Probability, Convex Optimization, ML, Deep Learning, RL, NLP, Speech Systems, AI For Humanity

WORK EXPERIENCE

ORACLE | DATA SCIENTIST 1

Bangalore, IN | Jul 2025 – Present

- Ship LLM query-understanding for Fusion Analytics Warehouse Copilot: entity classification, column reranking, and retrieval-backed generation.
- Improve Copilot quality via prompt tuning, JSON response contracts, and production backend-log analysis.

ORACLE | DATA ANALYTICS INTERN

Bangalore, IN | May 2024 – Jul 2024

- Built LLM synthetic-data pipeline for Oracle FAW (Python, K-Means/farthest-point sampling); adopted by a Consulting MTS.

PUBLICATIONS

Interpreting and Steering State-Space Models via Activation Subspace Bottlenecks | ICML 2026

Co-author; mechanistic interpretability on Mamba LLMs via activation subspace bottlenecks; benchmarked steering controllability vs. quality.

Stream-TTS: A Low-Latency Text-to-Speech using Kolmogorov-Arnold Networks | IEEE ICASSP 2025

Co-authored multilingual streaming TTS with KANs; 0.0795 RTF and 4.09 MOS with prosody control.

Fairness-Aware Submodular Optimization: A Survey | UNPUBLISHED SURVEY, IIT HYDERABAD

Sole-authored survey of fairness constraints in submodular maximization (summarization, influence spread); synthesized fairness-utility tradeoffs and open directions for equitable ML selection.

Attentive Multimodal Learning for Suicide Risk Assessment Using Acoustic and Prosodic Cues | UNDER REVIEW, SPEECH COMMUNICATION

Built initial multimodal pipeline with acoustic embeddings and prosodic cues for suicide risk classification.

RESEARCH ARTIFACTS

TINY-AYA GLOBAL BLINDSPOTS

PYTHON, HUGGINGFACE, LLM PROBING

Curated failure-case dataset (26/259) for CohereLabs/tiny-aya-global (3.35B); red-team probe across safety, bias, reasoning, and security scenarios.

ALLM FAIRNESS PROMPTS

PYTHON, HUGGINGFACE, LLM EVAL

Fairness evaluation prompt set for aligned language models; used to assess bias and stereotyping in multilingual LLM outputs.

PROJECTS

MULTILINGUAL LLM OVER-REFUSAL EVALUATION

PYTHON, PYTORCH, HUGGINGFACE, OR-BENCH

Replicated Sakonii task-over-refusal benchmark on Llama-3.1-8B with calibrated OR-Bench-style judge; extending evaluation to Hindi/Hinglish task wrappers with human-reviewed translations.

OPTICAL CHARACTER RECOGNITION FOR ANCIENT SANSKRIT MANUSCRIPTS

PYTHON, OCR, CRNN

OCR for degraded Sanskrit manuscripts via line segmentation, CRNN fine-tuning, and multimodal LLM evaluation.

AI-POWERED SESSION HIJACKING PREVENTION

PYTHON, XGBOOST, FLASK, REDIS

HTTP anomaly detector on CSIC 2010 (F1 0.97, SMOTE) with Redis-backed Flask inference and admin dashboard.

SUDOKU WEB APPLICATION

REACT, DJANGO, PYTHON

Full-stack Sudoku game with puzzle generation, backtracking solver, and interactive UI.

SPOTIDB

PYTHON, SPOTIFY API, MYSQL

Spotify-inspired music database modeling tracks, artists, albums, playlists, and user ratings via the Spotify API.

DATA STRUCTURES & ALGORITHMS FROM SCRATCH

C++

Core data structures and algorithms implemented from first principles with emphasis on memory management and complexity trade-offs.

SKILLS

Languages: Python, C++, C, SQL, JavaScript, Bash

Technology: PyTorch, HuggingFace, LLM Eval, Prompt Engineering, Retrieval/Reranking, XGBoost, Flask, Redis, Git, Docker, $\LaTeX$